

In [1]:

```
1 #Importing Libraries
```

In [5]:

```
1 import seaborn as sns # To use Seaborn functions
```

In [6]:

```
1 import matplotlib.pyplot as plt # This imports the pyplot module from the matplotlib
```

In [8]:

```
1 %matplotlib inline
```

In [9]:

```
1 import pandas as pd # In order to import and use the pandas Library
```

In [10]:

```
1 sns.set() #
```

In [11]:

```
1 import numpy as np #the code tells Python to bring the NumPy Library into your curr
```

In [12]:

```
1 from sklearn.preprocessing import StandardScaler # to transform input dataset valu
```

In [13]:

```
1 from sklearn.model_selection import train_test_split # To splitting data arrays into
```

In [14]:

```
1 from sklearn.metrics import classification_report, confusion_matrix #Compute conf
```

In [15]:

```
1 df = pd.read_csv('heart.csv') #T import dataset
```

In [16]:

```
1 df.head() #To show header of dataset
```

Out[16]:

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

In [17]:

```
1 df.info() #To show information of dataset
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trtbps      303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalachh    303 non-null   int64
8   exng        303 non-null   int64
9   oldpeak     303 non-null   float64
10  slp         303 non-null   int64
11  caa         303 non-null   int64
12  thall       303 non-null   int64
13  output      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

In [18]:

```
1 df.isnull()    #To check whether if any value is null
```

Out[18]:

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa
0	False	False	False	False	False	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False
...	...	...	...	...	...	...	...	...	...	...	...	...
298	False	False	False	False	False	False	False	False	False	False	False	False
299	False	False	False	False	False	False	False	False	False	False	False	False
300	False	False	False	False	False	False	False	False	False	False	False	False
301	False	False	False	False	False	False	False	False	False	False	False	False
302	False	False	False	False	False	False	False	False	False	False	False	False

303 rows × 14 columns

In [19]:

```
1 df.isnull().sum()    #It gives you pandas series of column names along with the sum of non-null values
```

Out[19]:

```
age      0
sex      0
cp       0
trtbps   0
chol     0
fbs      0
restecg  0
thalachh 0
exng     0
oldpeak  0
slp      0
caa      0
thall    0
output   0
dtype: int64
```

In [20]:

```
1 df_copy = df.copy(deep=True)    # To create a shallow copy of Pandas DataFrame
```

In [21]:

```
1 df_copy[['trtbps','chol','thalachh','oldpeak']] = df_copy[['trtbps','chol','thalachh','oldpeak']]
```

In [22]:

```
1 print(df_copy.isnull().sum()) #To find the total number of missing value
```

```
age      0
sex      0
cp       0
trtbps   0
chol     0
fbs      0
restecg   0
thalachh  0
exng     0
oldpeak   99
slp      0
caa      0
thall     0
output    0
dtype: int64
```

In [23]:

```
1 df_copy['oldpeak'].fillna(df_copy['oldpeak'].mean(), inplace = True) # To create a sh
```

In [24]:

```
1 color_wheel = {1: "#0392cf",2:"#7bc043"} #The short answer is determine the color of
```

In [25]:

```
1 colors = df["output"].map(lambda x:color_wheel.get(x + 1))
```


In [26]:

```
1 print(df.output.value_counts()) # to get the counts of unique values of the dataframe
```

```
1    165
0     138
Name: output, dtype: int64
```

In [27]:

```
1 p=df.output.value_counts().plot(kind="bar") # plot value_counts of Series
```




In [28]:

```
1 sns.heatmap(df.corr()) #To see the coreiation using heatmap
```

Out[28]:

&lt;AxesSubplot:&gt;



In [29]:

```
1 X = df.drop('output',axis=1)    # Remove the column from the DataFrame
2 print(X)
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	s1
p \											
0	63	1	3	145	233	1	0	150	0	2.3	
0											
1	37	1	2	130	250	0	1	187	0	3.5	
0											
2	41	0	1	130	204	0	0	172	0	1.4	
2											
3	56	1	1	120	236	0	1	178	0	0.8	
2											
4	57	0	0	120	354	0	1	163	1	0.6	
2											
..	...	...	..	...	...	...	...	...	...	...	
...											
298	57	0	0	140	241	0	1	123	1	0.2	
1											
299	45	1	3	110	264	0	1	132	0	1.2	
1											
300	68	1	0	144	193	1	1	141	0	3.4	
1											
301	57	1	0	130	131	0	1	115	1	1.2	
1											
302	57	0	1	130	236	0	0	174	0	0.0	
1											
	caa	thall									
0	0	1									
1	0	2									
2	0	2									
3	0	2									
4	0	2									
..	...	...									
298	0	3									
299	0	3									
300	2	3									
301	1	3									
302	1	2									

[303 rows x 13 columns]

In [30]:

```
1 y = df['output'] # show output after dropping
2 print(y)
```

```
0      1
1      1
2      1
3      1
4      1
..
298    0
299    0
300    0
301    0
302    0
```

Name: output, Length: 303, dtype: int64

In [31]:

```
1 #Splitting the data set into training and test data
2 from sklearn.model_selection import train_test_split, cross_val_score, cross_val_predict
```

In [32]:

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=7)
```

In [33]:

```
1 #Fitting Decision tree classifier to the training set
2 from sklearn.tree import DecisionTreeClassifier
```

In [34]:

```
1 dtree = DecisionTreeClassifier()
```

In [35]:

```
1 dtree.fit(X_train, y_train)
```

Out[35]:

DecisionTreeClassifier()

In [36]:

```
1 #Predicting the test set result
2 y_pred = dtree.predict(X_test)
```

In [37]:

```
1 print("Classification report - \n", classification_report(y_test,y_pred))
```

```
Classification report -
              precision    recall  f1-score   support

     0           0.84        0.53        0.65         30
     1           0.67        0.90        0.77         31

 accuracy          0.72         61
 macro avg         0.75         61
 weighted avg      0.75         61
```

In [38]:

```
1 #creating the confusion matrix
2 cm = confusion_matrix(y_test, y_pred)
```

In [39]:

```
1 plt.figure(figsize=(5,5))
```

Out[39]:

<Figure size 360x360 with 0 Axes>


<Figure size 360x360 with 0 Axes>

In [40]:

```
1 sns.heatmap(data=cm,linewidths= .5, annot=True, square=True, cmap = 'Blues')
2 plt.ylabel('Actual label')
3 plt.xlabel('Predicted label')
4 all_sample_title = 'Accuracy Score: {0}'.format(dtree.score(X_test, y_test))
5 plt.title(all_sample_title, size = 15)
6
```

Out[40]:

Text(0.5, 1.0, 'Accuracy Score: 0.7213114754098361')





In [41]:

```
1 from sklearn.metrics import roc_curve, roc_auc_score
```

In [42]:


```
1 y_pred_proba = dtree.predict_proba(X_test)[:][:,1]
```

In [43]:

```
1 df_actual_predicted=pd.concat([pd.DataFrame(np.array(y_test),columns=['y_actual']),pd
2 df_actual_predicted.index=y_test.index
3 fpr, tpr, tr = roc_curve(df_actual_predicted['y_actual'], df_actual_predicted['y_pred
4 auc = roc_auc_score(df_actual_predicted['y_actual'], df_actual_predicted['y_pred_prob
5 plt.plot(fpr,tpr,label='Auc=%0.4f'%auc)
6 plt.plot(fpr,fpr,linestyle='--',color='k')
7 plt.xlabel('False Positive Rate')
8 plt.ylabel('True Positive Rate')
9 plt.title('ROC Curve',size=15)
10 plt.legend()
11
12
```

Out[43]:

<matplotlib.legend.Legend at 0x271d318d700>



In [44]:

```
1 #fitting the svm classifier to the training set
2 from sklearn import svm
3 from sklearn.svm import SVC
4 svm = svm.SVC(kernel='linear',gamma='auto',probability=True)
5 svm.fit(X_train,y_train)
6 y_pred = svm.predict(X_test)
7 print("Classification report - \n", classification_report(y_test,y_pred))
8
9
10
11
12
```

```
Classification report -
              precision    recall  f1-score   support

     0       0.86        0.60       0.71         30
     1       0.70        0.90       0.79         31


 accuracy          0.75         61
 macro avg       0.78         0.75         0.75         61
weighted avg       0.78         0.75         0.75         61
```

In [45]:

```
1 cm = confusion_matrix(y_test, y_pred)
2 plt.figure(figsize=(5,5))
3 sns.heatmap(data=cm,linewidths= .5, annot=True, square=True, cmap = 'Blues')
4 plt.ylabel('Actual label')
5 plt.xlabel('Predicted label')
6 all_sample_title = 'Accuracy Score: {0}'.format(svm.score(X_test, y_test))
7 plt.title(all_sample_title, size = 15)
8
```

Out[45]:

Text(0.5, 1.0, 'Accuracy Score: 0.7540983606557377')




In [46]:

```
1 from sklearn.metrics import roc_curve, roc_auc_score
2 y_pred_proba = svm.predict_proba(X_test)[:][:,1]
3 df_actual_predicted=pd.concat([pd.DataFrame(np.array(y_test),columns=['y_actual']),pd
4 df_actual_predicted.index=y_test.index
5 fpr, tpr, tr = roc_curve(df_actual_predicted['y_actual'], df_actual_predicted['y_pred
6 auc = roc_auc_score(df_actual_predicted['y_actual'], df_actual_predicted['y_pred_proba
7 plt.plot(fpr,tpr,label='Auc=%0.4f'%auc)
8 plt.plot(fpr,fpr,linestyle='--',color='k')
9 plt.xlabel('False Positive Rate')
10 plt.ylabel('True Positive Rate')
11 plt.title('ROC Curve',size=15)
12 plt.legend()
13
14
```

Out[46]:

<matplotlib.legend.Legend at 0x271d1524e50>



In [47]:

```
1 #fitting logistic regression to the training set
2 from sklearn.linear_model import LogisticRegression
3 logreg=LogisticRegression()
4 logreg.fit(X_train ,y_train)
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:  
763: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

Out[47]:

LogisticRegression()

In [48]:

```
1 y_pred = logreg.predict(X_test)
2 print("Classification report - \n", classification_report(y_test,y_pred))
```

Classification report -


	precision	recall	f1-score	support
0	0.86	0.60	0.71	30
1	0.70	0.90	0.79	31
accuracy			0.75	61
macro avg	0.78	0.75	0.75	61
weighted avg	0.78	0.75	0.75	61

In [49]:

```
1 cm = confusion_matrix(y_test, y_pred)
2 plt.figure(figsize=(5,5))
3 sns.heatmap(data=cm,linewidths= .5, annot=True, square=True, cmap = 'Blues')
4 plt.ylabel('Actual label')
5 plt.xlabel('Predicted label')
6 all_sample_title = 'Accuracy Score: {0}'.format(logreg.score(X_test, y_test))
7 plt.title(all_sample_title, size = 15)
```

Out[49]:

Text(0.5, 1.0, 'Accuracy Score: 0.7540983606557377')




In [50]:

```
1 from sklearn.metrics import roc_curve, roc_auc_score
2 y_pred_proba = logreg.predict_proba(X_test)[:][:,1]
3 df_actual_predicted=pd.concat([pd.DataFrame(np.array(y_test),columns=['y_actual']),pd
4 df_actual_predicted.index=y_test.index
5 fpr, tpr, tr = roc_curve(df_actual_predicted['y_actual'], df_actual_predicted['y_pred
6 auc = roc_auc_score(df_actual_predicted['y_actual'], df_actual_predicted['y_pred_prob
7 plt.plot(fpr,tpr,label='Auc=%0.4f'%auc)
8 plt.plot(fpr,fpr,linestyle='--',color='k')
9 plt.xlabel('False Positive Rate')
10 plt.ylabel('True Positive Rate')
11 plt.title('ROC Curve',size=15)
12 plt.legend()
13
```

Out[50]:

<matplotlib.legend.Legend at 0x271d328f550>



In [51]:

```
1 #fitting K-NN classifier to the training data
2 from sklearn.neighbors import KNeighborsClassifier
3 knn=KNeighborsClassifier()
4 knn.fit(X_train,y_train)
5 y_pred = knn.predict(X_test)
6 print("Classification report - \n", classification_report(y_test,y_pred))
7
8
```

Classification report -


	precision	recall	f1-score	support
0	0.60	0.30	0.40	30
1	0.54	0.81	0.65	31
accuracy			0.56	61
macro avg	0.57	0.55	0.52	61
weighted avg	0.57	0.56	0.53	61

In [52]:

```
1 cm = confusion_matrix(y_test, y_pred)
2 plt.figure(figsize=(5,5))
3 sns.heatmap(data=cm,linewidths= .5, annot=True, square=True, cmap = 'Blues')
4 plt.ylabel('Actual label')
5 plt.xlabel('Predicted label')
6 all_sample_title = 'Accuracy Score: {0}'.format(knn.score(X_test, y_test))
7 plt.title(all_sample_title, size = 15)
```

Out[52]:

Text(0.5, 1.0, 'Accuracy Score: 0.5573770491803278')



In [53]:

```
1 from sklearn.metrics import roc_curve, roc_auc_score
2 y_pred_proba = knn.predict_proba(X_test)[:][:,1]
3 df_actual_predicted=pd.concat([pd.DataFrame(np.array(y_test),columns=['y_actual']),pd
4 df_actual_predicted.index=y_test.index
5 fpr, tpr, tr = roc_curve(df_actual_predicted['y_actual'], df_actual_predicted['y_pred
6 auc = roc_auc_score(df_actual_predicted['y_actual'], df_actual_predicted['y_pred_prob
7 plt.plot(fpr,tpr,label='Auc=%0.4f'%auc)
8 plt.plot(fpr,fpr,linestyle='--',color='k')
9 plt.xlabel('False Positive Rate')
10 plt.ylabel('True Positive Rate')
11 plt.title('ROC Curve',size=15)
12 plt.legend()
```

Out[53]:

<matplotlib.legend.Legend at 0x271d33c5dc0>

